


DIMENSIONES DE THREADS Y BLOCKS



gridDim - indica el número de bloques en un grid.
blockDim - indica el número de threads en un bloque.
blockIdx - índice del bloque dentro del grid.
threadIdx - índice del thread dentro del bloque.

- Los bloques puede trabajar en 1D ó 2D
- Los threads pueden trabajar en 1D, 2D ó 3D

dim3

- **dim3** - es un tipo de vector entero usado en cuda para indicar las dimensiones en el grid o bloque cuando se invoca el kernel.

- **dim3** puede tomar 1, 2 ó 3 argumentos

```
dim3 varDim1D(x);
```

```
dim3 varDim2D(x, y);
```

```
dim3 varDim3D(x, y, z);
```

Actividad 4.1

- Realizar una función que sume dos matrices, pero los argumentos deben ser apuntadores sencillos.
- Construir el main para comprobar la función.
- Corroborar el casting adecuado.

Sumar dos matrices en CUDA

- Se van a sumar cada uno de los elementos de las matrices a sumar en forma simultánea

$$A = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}$$

$$B = \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

$$C = \begin{bmatrix} C_{11}=A_{11}+B_{11} & C_{12}=A_{12}+B_{12} \\ C_{21}=A_{21}+B_{21} & C_{22}=A_{22}+B_{22} \end{bmatrix}$$

$$C_{11} = A_{11} + B_{11}$$

Procesador 1

$$C_{12} = A_{12} + B_{12}$$

Procesador 2

$$C_{21} = A_{21} + B_{21}$$

Procesador 3


$$C_{22} = A_{22} + B_{22}$$

Procesador 4

```
1  #include<stdio.h>
2  #define F 2
3  #define C 3
4  __global__ void sumaMat(int a[], int b[], int c[]) {
5      int x = threadIdx.x; //Filas
6      int y = threadIdx.y; //Columnas
7      int tid = y + (C*x);
8      c[tid] = a[tid] + b[tid];
9  }
10 int main() {
11     int numBlocks = 1;
12     int i, j;
13     dim3 threadsPerBlock(F, C);
14     int a_h[F][C]={ {1,1,1}, {3,3,3} };
15     int b_h[F][C]={ {2,2,2}, {4,4,4} };
16     int c_h[F][C];
17     int *a_d, *b_d, *c_d;
18     int size = F*C*sizeof(int);
```

```
19  cudaMalloc((void **)&a_d, size);
20  cudaMalloc((void **)&b_d, size);
21  cudaMalloc((void **)&c_d, size);
22  cudaMemcpy(a_d, a_h, size, cudaMemcpyHostToDevice);
23  cudaMemcpy(b_d, b_h, size, cudaMemcpyHostToDevice);
24  sumaMat<<<numBlocks, threadsPerBlock>>>(a_d, b_d, c_d);
25  cudaMemcpy(c_h, c_d, size, cudaMemcpyDeviceToHost);
26
27  for(i = 0; i < F; i++)
28      for(j = 0; j < C; j++)
29          printf("c_h[%d][%d]=%d\n",i, j, c_h[i][j]);
30
31  cudaFree(a_d);
32  cudaFree(b_d);
33  cudaFree(c_d);
34  return 0;
35 }
```

```
#define F 2
#define C 3
...
int numBlocks = 1;
dim3 threadsPerBlock(F, C);
...
sumaMat<<<numBlocks, threadsPerBlock
    >>>(a_d, b_d, c_d);
...
```




```
int numBlocks = 1;
dim3 threadsPerBlock(2, 3);
sumaMat<<<numBlocks, threadsPerBlock>>>(...);
gridDim.x      = 1
blockDim.x      = 2
blockDim.y      = 3
threadIdx.x     = 0, 1
threadIdx.y     = 0, 1, 2
```

Actividad 4.2

- Imprimir los valores anteriores utilizando la biblioteca `cuPrintf.cu`, la salida seria:

```
$ ./imprimeDatosMat
```

```
[0, 0]: gd=1 bdx=2 bdy=3 tx=0 ty=0
```

```
[0, 1]: gd=1 bdx=2 bdy=3 tx=1 ty=0
```

```
[0, 2]: gd=1 bdx=2 bdy=3 tx=0 ty=1
```

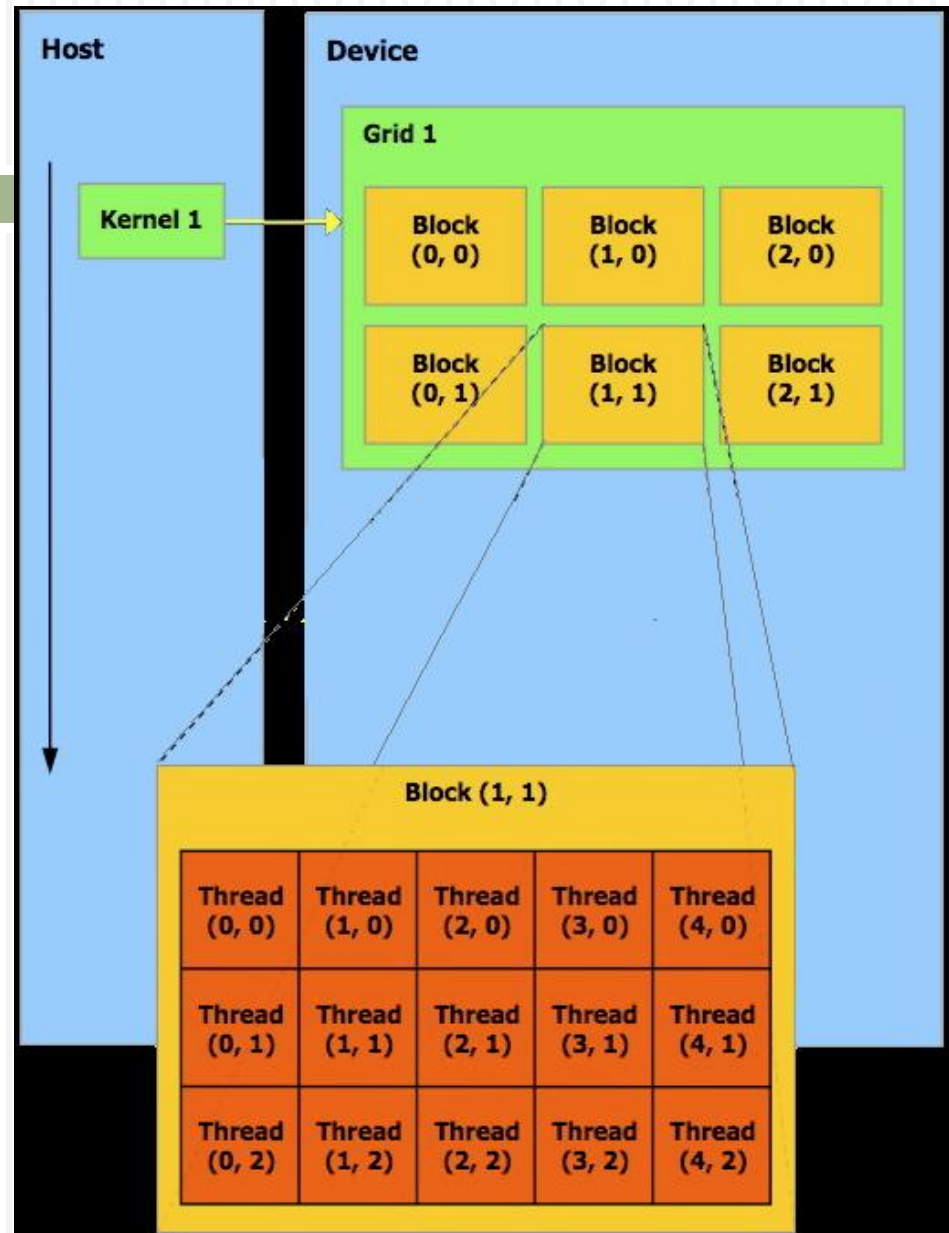
```
[0, 3]: gd=1 bdx=2 bdy=3 tx=1 ty=1
```

```
[0, 4]: gd=1 bdx=2 bdy=3 tx=0 ty=2
```

```
[0, 5]: gd=1 bdx=2 bdy=3 tx=1 ty=2
```

Ejercicio

- Calcular la suma de dos matrices de 6×15 .
- Con un grid de 2×3 bloques.
- Con bloques de 3×5 threads.
- Se maneja la sintaxis Filas x Columnas.



Código del kernel

```
__global__ void sumaMat(int a[], int b[], int c[])
{
    int x = threadIdx.x + blockDim.x*blockIdx.x; //Fil
    int y = threadIdx.y + blockDim.y*blockIdx.y; //Col
    ...
}
```

Llamada al kernel

```
...
dim3 numBlocks(2, 3);           //Grid structure
dim3 threadsPerBlock(3, 5);     //Block structure
...
//Invocando al kernel
sumaMat<<<numBlocks, threadsPerBlock>>>(a_d, b_d, c_d);
...
```

```
dim3 numBlocks(2, 3);  
dim3 threadsPerBlock(3, 5);  
sumaMat<<<numBlocks, threadsPerBlock>>>(...);  
gridDim.x      = 2  
gridDim.y      = 3  
blockDim.x     = 3  
blockDim.y     = 5  
blockIdx.x     = 0, 1  
blockIdx.y     = 0, 1, 2  
threadIdx.x    = 0, 1, 2  
threadIdx.y    = 0, 1, 2, 3, 4
```

Actividad 4.3

- Imprimir los valores para el ID del bloque y de los threads, utilizando la biblioteca `cuPrintf.cu`, la salida seria:

```
$/imprimeDatosMat_2
```

```
[0, 0]: bx =0 by =0 tx =0 ty =0
```

```
[0, 1]: bx =0 by =0 tx =1 ty =0
```

```
[0, 2]: bx =0 by =0 tx =2 ty =0
```

```
[0, 3]: bx =0 by =0 tx =0 ty =1
```

```
[0, 4]: bx =0 by =0 tx =1 ty =1
```

```
[0, 5]: bx =0 by =0 tx =2 ty =1
```

[0, 6]: $bx = 0$ $by = 0$ $tx = 0$ $ty = 2$

[0, 7]: $bx = 0$ $by = 0$ $tx = 1$ $ty = 2$

[0, 8]: $bx = 0$ $by = 0$ $tx = 2$ $ty = 2$

[0, 9]: $bx = 0$ $by = 0$ $tx = 0$ $ty = 3$

[0, 10]: $bx = 0$ $by = 0$ $tx = 1$ $ty = 3$

[0, 11]: $bx = 0$ $by = 0$ $tx = 2$ $ty = 3$

[0, 12]: $bx = 0$ $by = 0$ $tx = 0$ $ty = 4$

[0, 13]: $bx = 0$ $by = 0$ $tx = 1$ $ty = 4$

[0, 14]: $bx = 0$ $by = 0$ $tx = 2$ $ty = 4$

[1, 0]: $bx = 1$ $by = 0$ $tx = 0$ $ty = 0$

[1, 1]: $bx = 1$ $by = 0$ $tx = 1$ $ty = 0$

...

Referencias

- Sito de NVIDIA, <https://developer.nvidia.com/>
- CUDA C PROGRAMMING GUIDE, NVIDIA
- CUDA by Examples, NVIDIA